

Making HTTP **realtime** with HTTP 2.0

aka, dropping the hacks, reclaiming performance!

Ilya Grigorik

igrigorik@google.com

Google

```
$> telnet igvita.com 80
```

```
Connected to 173.230.151.99
```

```
GET /archive
```

```
Hypertext delivery with HTTP 0.9! - eom.  
(connection closed)
```

HTTP 0.9 is the ultimate MVP - one line, plain-text
“protocol” to test drive the “www idea”.



```
$> telnet ietf.org 80
```

```
Connected to 74.125.xxx.xxx
```

```
GET /rfc/rfc1945.txt HTTP/1.0
```

```
User-Agent: CERN-LineMode/2.15 libwww/2.17b3
```

```
Accept: */*
```

```
HTTP/1.0 200 OK
```

```
Content-Type: text/plain
```

```
Content-Length: 137582
```

```
Last-Modified: Wed, 1 May 1996 12:45:26 GMT
```

```
Server: Apache 0.84
```

```
4 years of rapid iteration later... eom.
```

```
(connection closed)
```

HTTP 1.0 is an informational RFC - documents
“common usage” of HTTP found in the wild.



```
$> telnet google.com 80
```

```
Connected to 74.125.xxx.xxx
```

```
GET /index.html HTTP/1.1
```

```
Host: website.org
```

```
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_7_4)... (snip)
```

```
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,*/*;q=0.8
```

```
Accept-Encoding: gzip,deflate,sdch
```

```
Accept-Language: en-US,en;q=0.8
```

```
Cookie: __qca=P0-800083390... (snip)
```

```
HTTP/1.1 200 OK
```

```
Connection: keep-alive
```

```
Transfer-Encoding: chunked
```

```
Server: nginx/1.0.11
```

```
Content-Type: text/html; charset=utf-8
```

```
Date: Wed, 25 Jul 2012 20:23:35 GMT
```

```
Expires: Wed, 25 Jul 2012 20:23:35 GMT
```

```
Cache-Control: max-age=0, no-cache
```

```
100
```

```
<!doctype html>
```


```
(snip)
```

***HTTP 1.1 ships as RFC standard in 1999 - hyper
{text}media all the things!***



15 years later!

HTTPbis is getting dangerously close to closing all the outstanding 1.1 bugs. :)
In the meantime a few things have happened since...



Geocities ftw!
(circa 1997-2000)



Then AJAX happened.
(circa 2002-2005)



WebGL
WebRTC
Web Audio
WebSockets
...
Mobile web


(today)



State of the **HTTP nation**...

- **11** distinct hosts per page
- **76** distinct requests per page
- **987** transferred KBytes per page

- **Concurrency: 1**
- **Parallelism: 6**



*One transfer per connection.
At most six parallel connections.*


*Resulting in typical load times of **3.8-10.5 seconds**.*



* All numbers are medians, based on [HTTP Archive crawl data](#).

@igrigorik

For all its awesome, **HTTP 1.X has a lot of perf problems...**



Top Desktop			
name	score	PerfTiming	Connections per Hostname
<input type="checkbox"/> Chrome 20 →	12/16	yes	6
<input type="checkbox"/> Firefox 14 →	13/16	yes	6
<input type="checkbox"/> IE 8 →	7/16	no	6
<input type="checkbox"/> IE 9 →	12/16	yes	6
<input type="checkbox"/> Opera 12 →	10/16	no	6
<input type="checkbox"/> RockMelt 0.9 →	13/16	yes	6
<input type="checkbox"/> Safari 5.1 →	12/16	no	6

- Limited parallelism (~6)
- Client-side request queuing
- Browser queueing heuristics
- High protocol overhead
 - ~800 bytes + cookies
- Competing TCP flows
- Spurious retransmissions
- Extra memory / FD resources
- Handshake overhead
- Slow-start overhead
- ...



Where there's a will, there's a way...

we're an inventive bunch, so we came up with some "optimizations" (read, "hacks")



- **Domain sharding**

- TCP Slow Start? Browser limits, Nah... 15+ parallel requests -- *Yeehaw!!!*
- Causes congestion and unnecessary latency and retransmissions

- **Concatenating files (JavaScript, CSS)**

- Reduces number of downloads and latency overhead
- Less modular code and expensive cache invalidations (e.g. app.js)
- Slower execution (must wait for entire file to arrive)

- **Spriting images**

- Reduces number of downloads and latency overhead
- Painful and annoying preprocessing and expensive cache invalidations
- Have to decode entire sprite bitmap - CPU time and memory

- **Resource inlining**

- Eliminates the request for small resources
- Resource can't be cached, inflates parent document
- 30% overhead on base64 encoding





Ok, that sucked. Let's **fix HTTP** instead...



*"**HTTP 2.0** is a protocol designed for **low-latency transport of content** over the World Wide Web"*

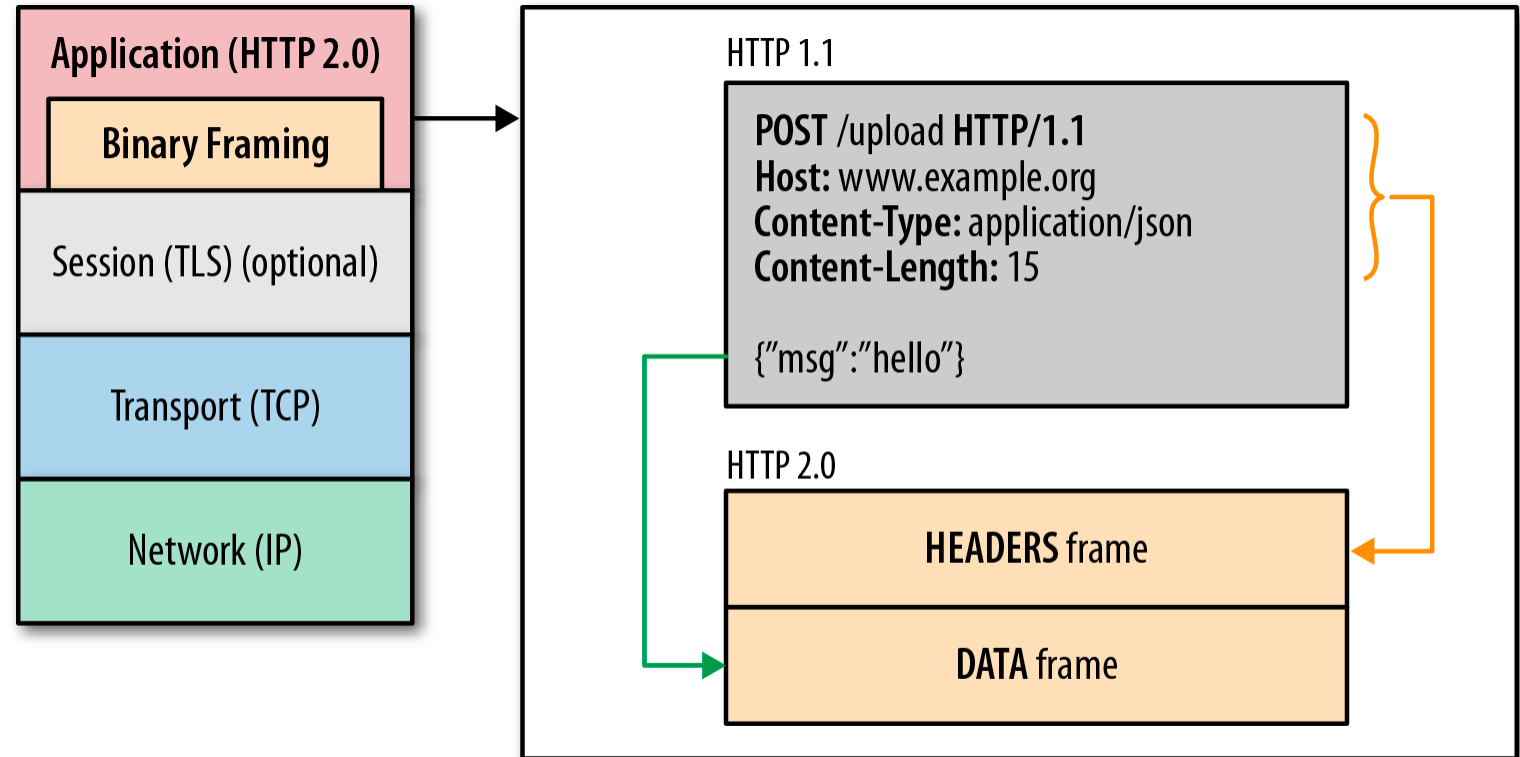
- Improve end-user perceived latency
- Address the "head of line blocking"
- Not require multiple connections
- Retain the semantics of HTTP/1.1

HTTP 2.0 work kicked off in Jan 2012.



HTTP 2.0 in one slide...

- **One TCP connection**
- **Request = Stream**
 - Streams are multiplexed
 - Streams are prioritized
- **Binary framing layer**
 - Prioritization
 - Flow control
 - Server push
- **Header compression**



*“... **we’re not replacing all of HTTP** — the methods, status codes, and most of the headers you use today will be the same. Instead, **we’re redefining how it gets used “on the wire” so it’s more efficient**, and so that it is more gentle to the Internet itself”*

- Mark Nottingham (HTTPbis chair)



Binary framing crash course in one slide...

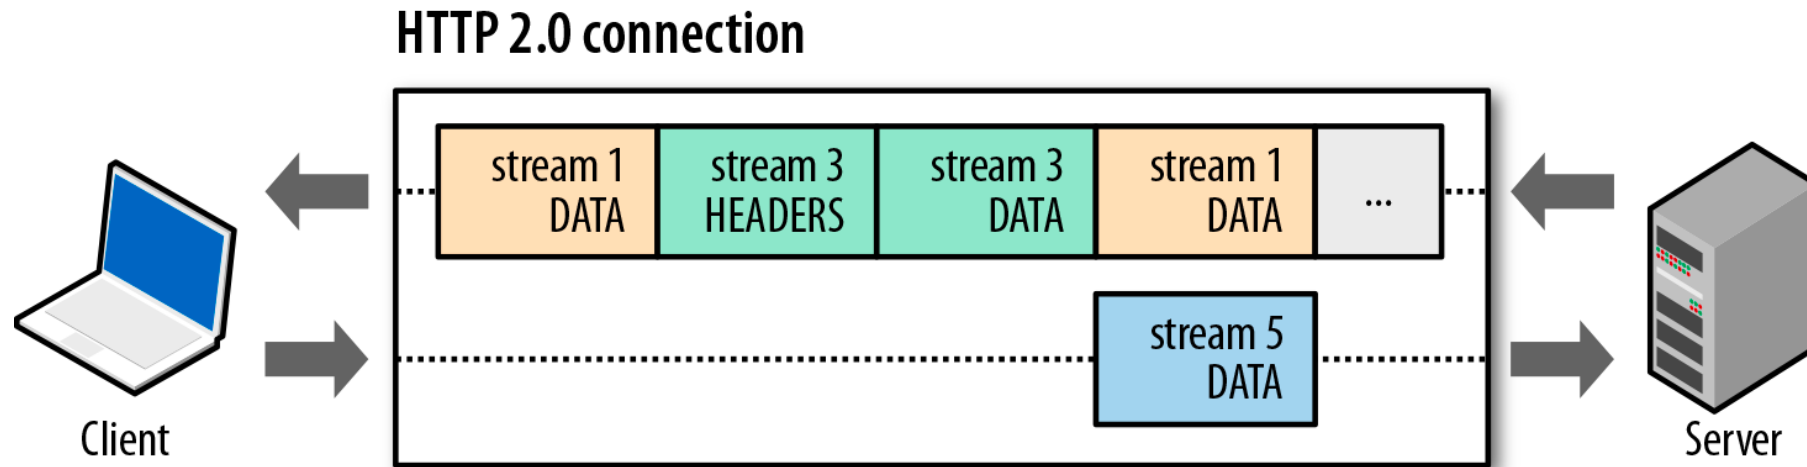
Bit	+0..7		+8..15	+16..23	+24..31
0	Length			Type	Flags
32	R	Stream Identifier			
...	Frame Payload				

- Length-prefixed frames
- All frames have same 8-byte header
- **Type** indicates ... type of frame:
 - *DATA, HEADERS, PRIORITY, PUSH_PROMISE, ...*
- Each frame may have custom **flags**
 - e.g. *END_STREAM*
- Each frame carries a **31-bit stream identifier**
 - After that, it's frame specific payload

```
frame = buffer.read(8)
if frame_i_care_about
    do_something_smart
else
    buffer.skip(frame.length)
end
```



Basic data flow in HTTP 2.0...




- **Streams are multiplexed** by splitting communication into frames
 - e.g. HEADERS, DATA, etc.
- **Frames are interleaved**
 - Frames can be prioritized
 - Frames can be flow controlled

*E.g. Please send **critical.css** with **priority 1**, please! **kittens.jpg** with **priority 10**.*



Connection + stream **flow-control**!



Stream flow-control enables **fine-grained resource control** between streams. E.g...

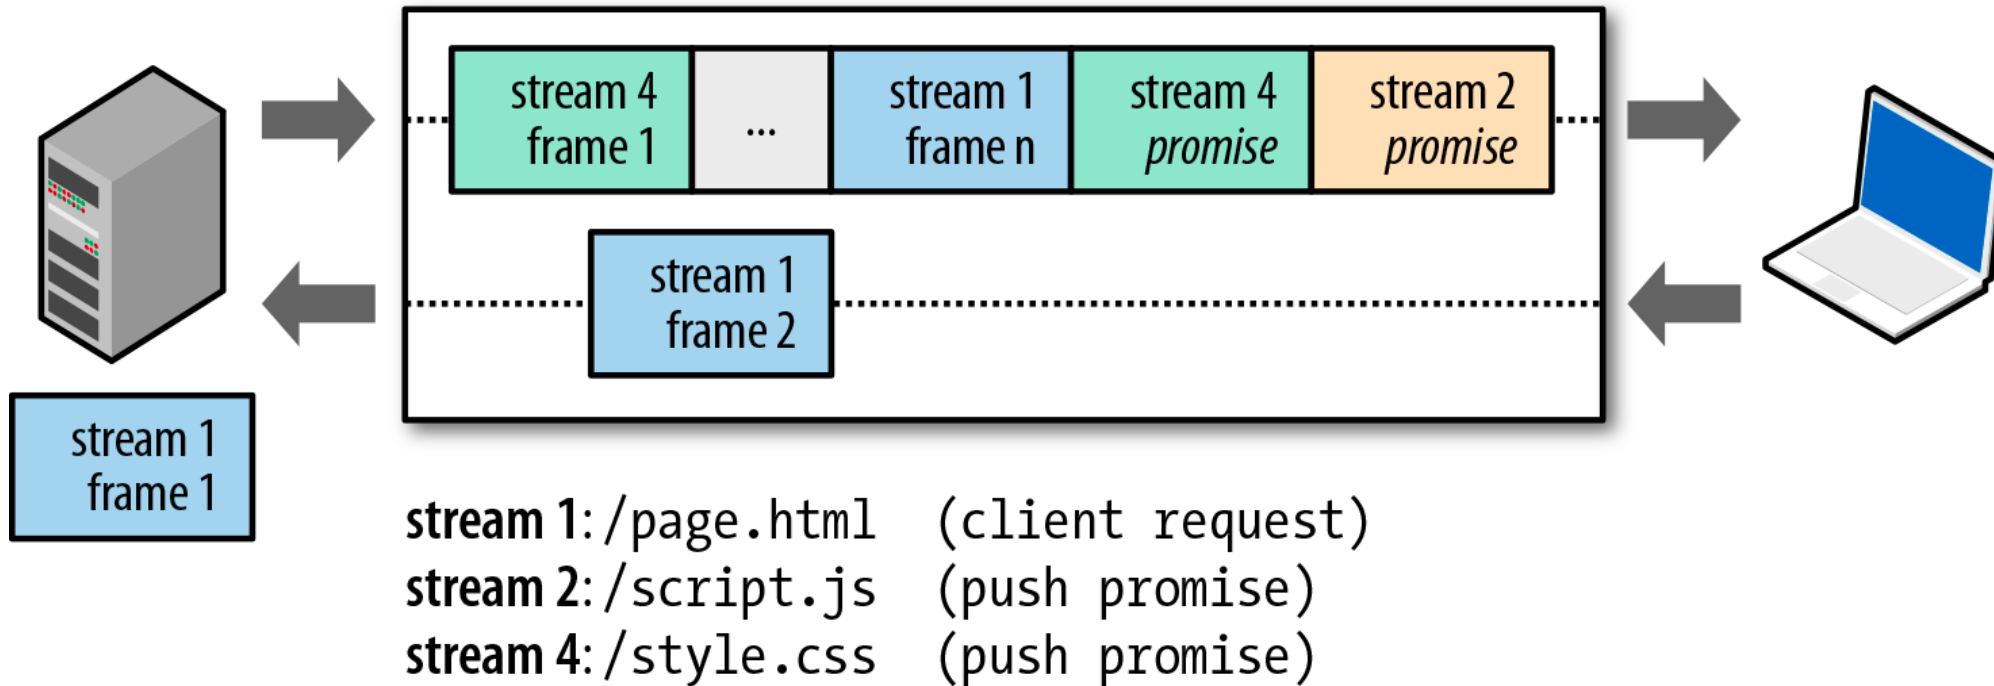
- **T(0)**: I am willing to receive **64 KB** of kittens.jpg.
- **T(0)**: I am willing to receive **500KB** of critical.js
- ...
- **T(n)**: Ok, now send the **remainder** of kittens.jpg.

Client controls how and when the stream and connection window is incremented!



Server push... is replacing inlining.

HTTP 2.0 connection



*Inlining **is** server push. Except, **HTTP 2.0 server push is cacheable!***

- One request, multiple responses.
- Push redirects, cache invalidations, ... lots of new and exciting possibilities!



HTTP 2.0 provides **header compression!**

Request #1

:method	GET
:scheme	https
:host	example.com
:path	/resource
accept	image/jpeg
user-agent	Mozilla/5.0 ...



HEADERS frame (Stream 1)


```
:method: GET
:scheme: https
: host: example.com
: path: /resource
accept: image/jpeg
user-agent: Mozilla/5.0 ...
```

- Both sides maintain “header tables”
- New requests “toggle” or “insert” new values into the table
- New header set is a “diff” of the previous set of headers
- Repeat request (polling) with exact same headers incurs **no overhead**

*Byte cost of new stream: **8 bytes!** **



** as low as 8 bytes for an identical request.*



tl;dr: framing, mux, prioritization, flow control

same HTTP protocol semantics, complete performance reboot!



Benefits

- $\min(\text{request overhead}) = 8 \text{ bytes}$
- $\max(\text{parallelism}) = 100 \sim 1000 \text{ streams}$
- $\max(\text{client queueing latency}) = 0 \text{ ms}$
- Stream multiplexing
- Stream prioritization
- Stream flow control
- Server push
- **Plus all the usual HTTP goodies!**
 - Browser caching, authentication, content negotiation, content-encoding, transfer encoding, encryption, ...

*Available in your browser circa ~2014! **



** or today, via SPDY.*

Opportunities

- **Unshard, un-concat, un-sprite, ...**
 - Modular code, modular assets
 - Improved cache performance
 - Improved page load times
- **Layer other browser transports**
 - XHR / SSE: work out of the gate
 - (TODO) WebSockets over HTTP 2.0: free mux, flow control, prioritization
- **We need smarter servers!**
 - Support for prioritization, flow control, and resource push!
- **One RPC stack to rule them all**
 - Standardized, high-performance RPC protocol stack (internal, external)

*<crazy idea> **HTTP 2.0 over UDP? QUIC!** </crazy idea> **

** we'll leave that for RealTimeConf 2014 ;-)*

```
igrigorik { /git/http-2 } > ruby example/client.rb https://twitter.com/
NPN protocols supported by server: ["HTTP-draft-06/2.0", "spdy/3.1", "spdy/3", "http/1.1"]
Sending HTTP 2.0 request
Sending bytes: "PRI * HTTP/2.0\r\n\r\nSM\r\n\r\n"
Sending bytes: "\x00\x10\x04\x00\x00\x00\x00\x00\x00\x00\x04\x00\x00\x00d\x00\x00\x00a\x00\x00\xFF\xFF"
Sending bytes: "\x00\x19\x01\x05\x00\x00\x00\x01\x81\x84C\x0Ftwitter.com:443\x83F\x03*/*"
[Stream 1]: closing client-end of the stream
[Stream 1]: response headers: {":status"=>"200 OK", "cache-control"=>"no-cache, no-store, must-revalidate, pre-check=
date"=>"Wed, 09 Oct 2013 22:45:01 GMT", "expires"=>"Tue, 31 Mar 1981 05:00:00 GMT", "last-modified"=>"Wed, 09 Oct 2
d=v1%3A138135870145166583; Domain=.twitter.com; Path=/; Expires=Fri, 09-Oct-2015 22:45:01 UTC", "status"=>"200 OK",
saction"=>"ee18608b67c977ba", "x-ua-compatible"=>"IE=10,chrome=1", "x-xss-protection"=>"1; mode=block"}
[Stream 1]: response data chunk: <<<!DOCTYPE html>
<!--[if IE 8]><html class="lt-ie10 ie8 " lang="en"
```

github.com/igrigorik/http-2

*HTTP 2.0 draft 6 implementations in Firefox, Chrome, IE11
has SPDY support (stepping stone). Plus client/server
implementations in C, JavaScript (node.js), Ruby, Perl, etc.*

HTTP 2.0 is coming to a client / server near you in 2014.





HTTP 2.0: hpbn.co/http2



Ilya Grigorik

igrigorik@google.com

@igrigorik